

CachePool: Many-core cluster of customizable, lightweight scalar-vector PEs for irregular L2 data-plane workloads

Integrated Systems Laboratory (ETH Zürich)

Zexin Fu, Diyou Shen


zexifu, dishen@iis.ee.ethz.ch

Alessandro Vanelli-Coralli
Luca Benini

avanelli@iis.ee.ethz.ch
lbenini@iis.ee.ethz.ch

PULP Platform

Open Source Hardware, the way it should be!



@pulp_platform



pulp-platform.org



youtube.com/pulp_platform



Outline

- **Recall: The target workload & Extracted kernels**
- **Initially proposed architecture candidates**
 - Private L1 + Shared L2
 - Share L1 with configurable partition
- **Open discussions**



Recall: The Target Workload



- **RLC Packet Handling**

- Massive **unstructured sparse data handling** in double **linked-list** format -> **Data Cache**
- **Minimum** calculation overhead
- Need **buffer** for ACK retransmission handling

-> **Control-oriented**

-> **Large memory management**

- **RLC Control**

- Mixed scalar/vector instruction (40%-60%)
- Need to support different types of INT (8, 16, 32, 64)

-> **Need vector unit**

-> **64-bit data path**

- **Performance target**

- High throughput
- Tight TTI

-> **Manycore cluster**

-> **Low latency interconnect**



Recall: Extracted Kernels



- **RLC Packet Handling**

- Linked List
- Pointer Chasing

- **RLC Control**

- Sparse Matrix-Vector Multiplication(SpMV)
- Maximum Value Sorting
- Logarithm Calculation Using Taylor Expansion
- Sum Reduction



Tiles with Scalar+Vector cores,
Large size of data cache



Overhead of HW Managed Coherence



- **Snoop-based protocol**
 - Coherence traffic in the system interconnect
 - Long snoop latency
- **Directory-based protocol**
 - Extra area and power for directory structure
 - The directory can be the bandwidth bottleneck
- **Ways to avoid the coherence overhead**
 - Write-through cache
 - Shared cache
 - Software data management
- **Our initial candidate designs avoid using invalidate-based coherence protocol**
 - We are still open to any low-cost scalable coherence design



Design Candidate 1: Private Write-through L1 + Shared L2

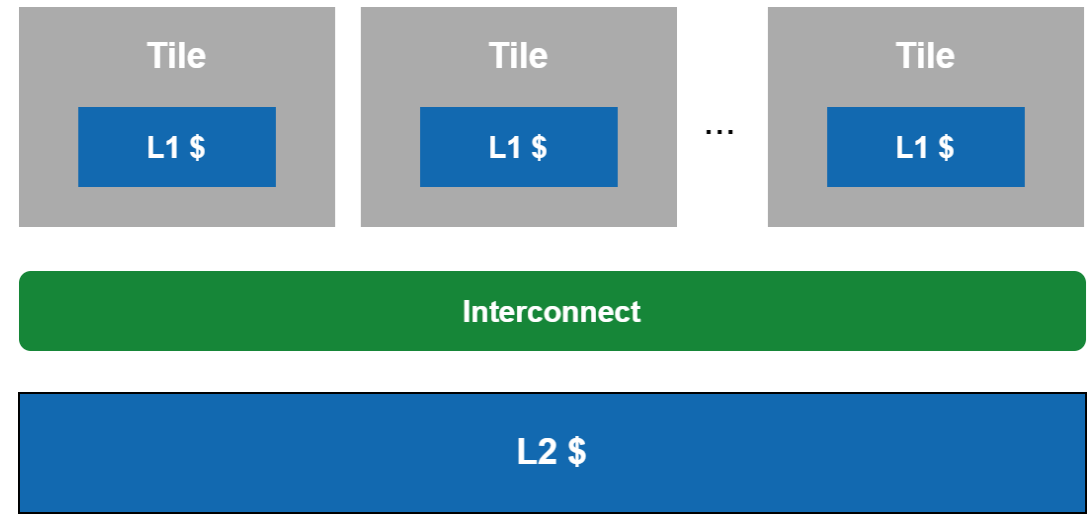


- **Small inclusive private L1 cache**

- To handle frequently accessed data
- Write-through to avoid cache coherence data transfer overhead
- Very small to reduce the data duplication

- **Large shared L2 cache**

- PE direct write through to L2
- Multi-sliced for large bandwidth



The Cluster-Level System Diagram



Private Write-through L1 + Shared L2



- **Pros**

- Higher data locality within a Tile.
- Reduce the interconnect cost between Tiles.

- **Cons**

- Data duplication
 - Multiple copies in different Private L1, and between L1 and L2, influence the actual usable memory size
- Write-through
 - The latency and bandwidth for writes. Performance and power degrade for write-intensive kernels.



Shared L1 Cache



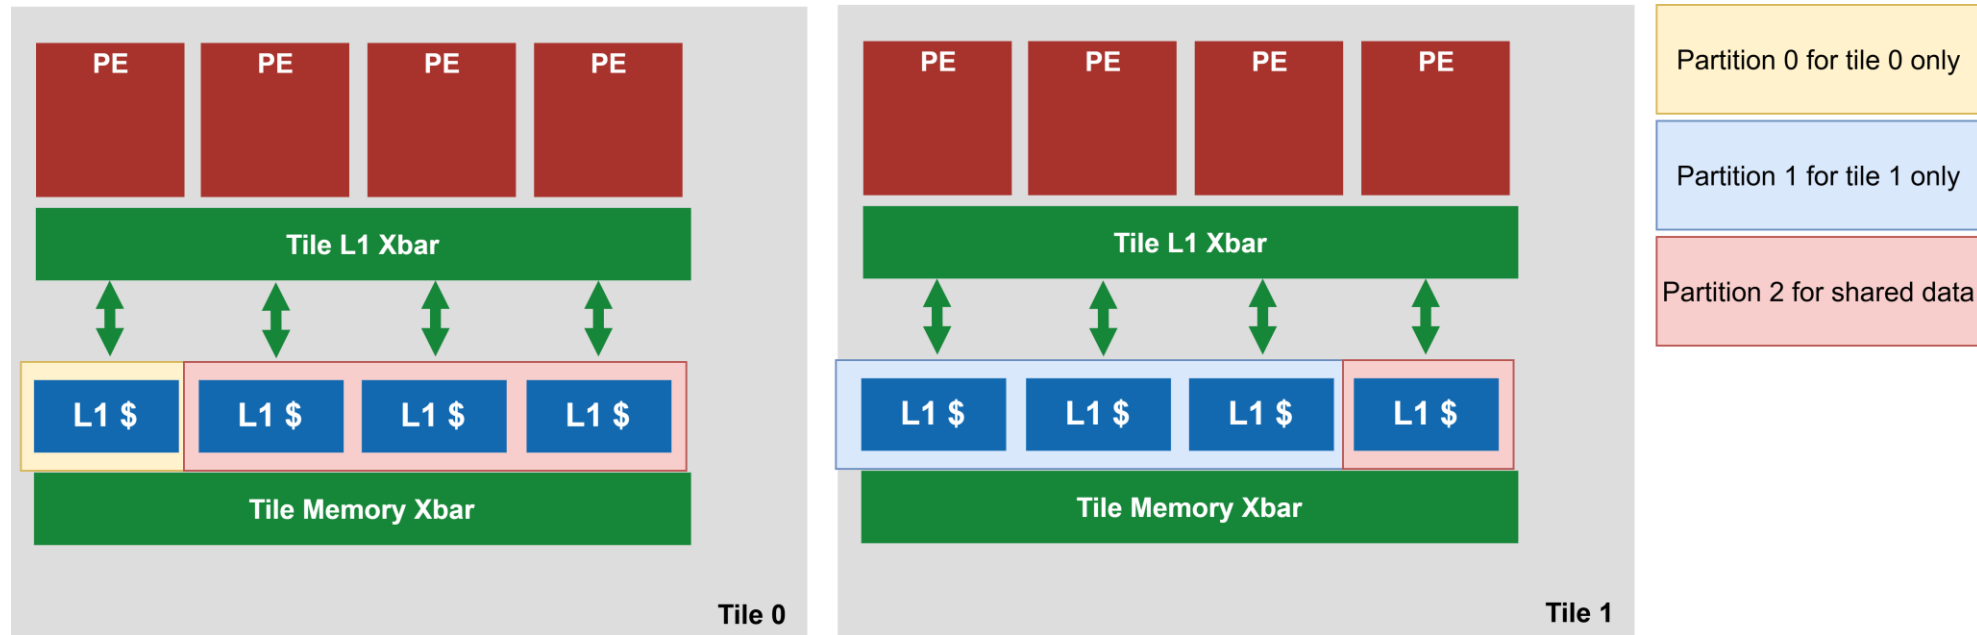
- **Traditional: Shared L1 Cache Design General View**
 - Similar to TeraPool/MemPool interconnection hierarchy.
 - Replace SPM with Cache
 - A large multi-banked L1 cache shared by all Tiles.
 - Each Tile will have its own Cache Controllers to manage the cache slices (multi-banked).
- **Pros**
 - No coherence policy is needed
 - No overhead from L2: Duplicated data, Tag, Control logic
- **Cons**
 - The physical-feasible interco. will limit the remote accessing BW
 - High ratio of long-latency remote access



Design Candidate 2: Shared L1 Cache with Partitioning



- Add partition on top of the shared multi-banked L1.
 - Configurable partition schemes for Tile-private or global-shared use.
 - The granularity of partition would be the cache bank.
 - Keep a globally shared partition to handle shared data and coherence.
- This can make keep private data closer to the consumer PEs to increase locality.



Shared L1 Cache with Partitioning



- **Pros**

- Same as previous, no coherence and no L2 overhead.
- Higher data locality, less remote access.
- Flexible cache resource allocation, maximizes the effective cache capacity utilization.

- **Cons**

- Configurable partitioning logic overhead in timing, area and design effort



Open Discussions



- **Design choices**
 - Private L1 + Shared L2
 - Share L1 with configurable partition
 - The design is not settled, open to any new design choices.
- **Kernel information**
 - Kernel parameters sheet and document mentioned in the last meeting
- **Next meeting schedule**



Thank you!

Q&A

