

CachePool: Many-core cluster of customizable, lightweight scalar-vector PEs for irregular L2 data-plane workloads

Integrated Systems Laboratory (ETH Zürich)

Zexin Fu, Diyou Shen

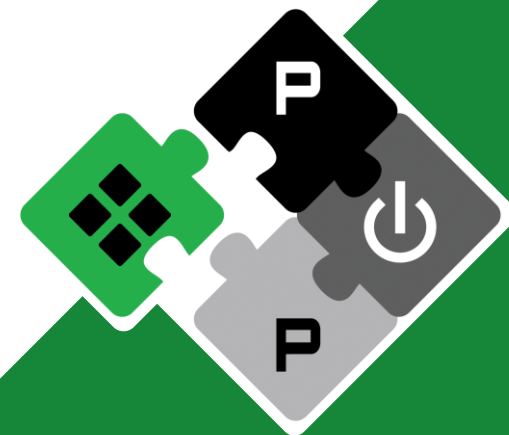
zexifu, dishen@iis.ee.ethz.ch

Alessandro Vanelli-Coralli
Luca Benini

avanelli@iis.ee.ethz.ch
lbenini@iis.ee.ethz.ch

PULP Platform

Open Source Hardware, the way it should be!



@pulp_platform



pulp-platform.org



youtube.com/pulp_platform



Status Update: Hardware



- **Finish the implementation of a single Tile**

- Four 64b Snitch-Spatz4 PEs
- Four 256b-wide cache banks
- Programmable selection in XBar
 - Can apply different mux offset bits from a CSR
 - Prepare for the partition functionality




Status Update: Hardware



- **Finish the implementation of a single Tile**

- Four 64b Snitch-Spatz4 PEs
- Four 256b-wide cache banks
- Programmable selection in XBar

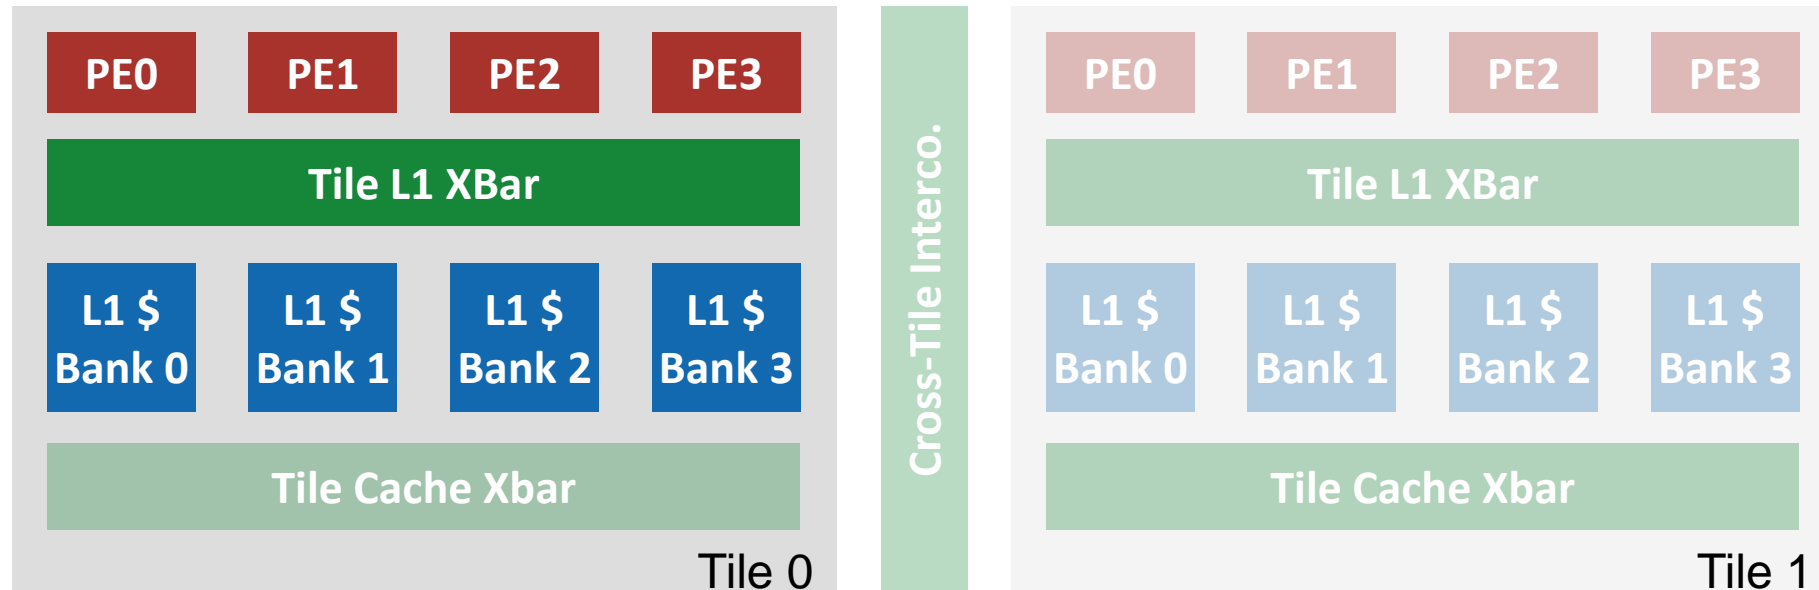


Status Update: Hardware



- **Finish the implementation of a single Tile**

- Four 64b Snitch-Spatz4 PEs
- Four 256b-wide cache banks
- Programmable selection in XBar

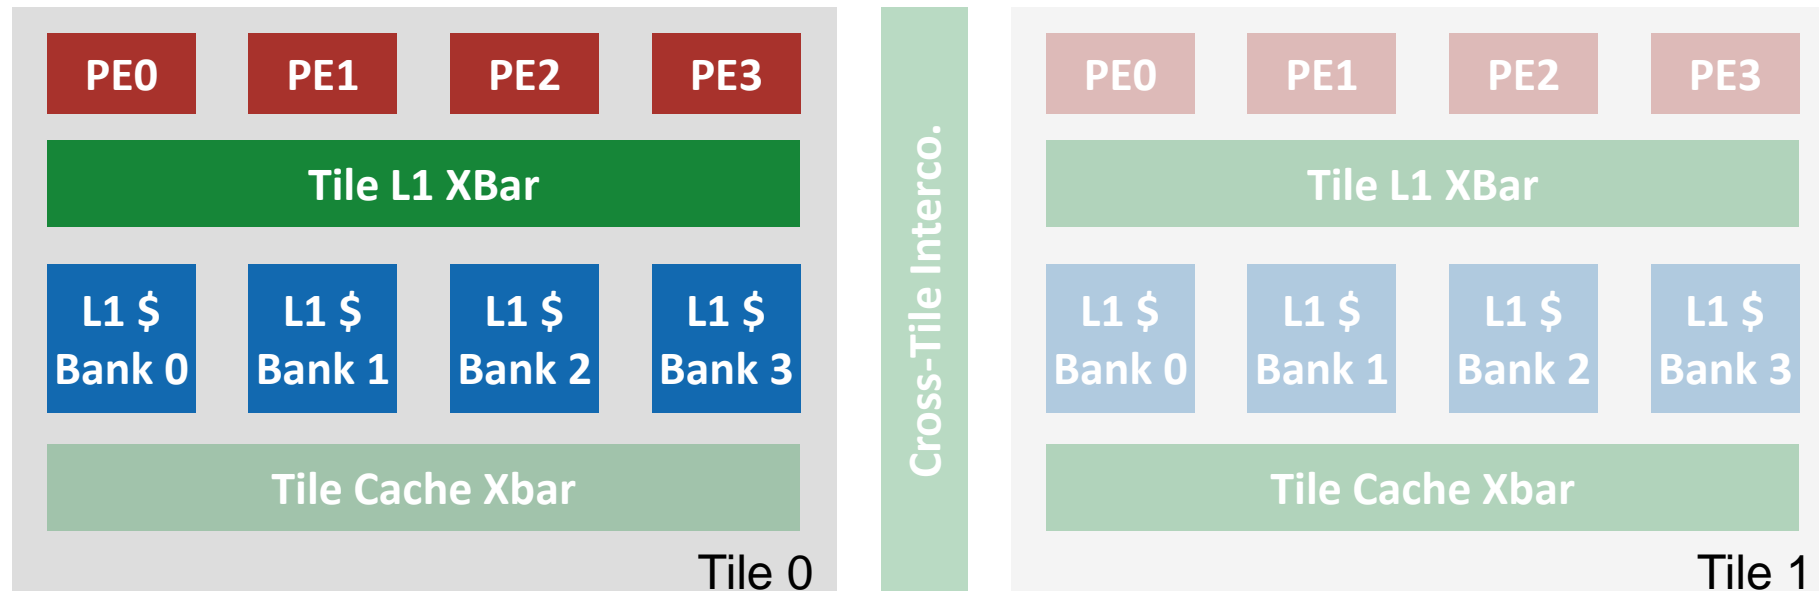


Status Update: Hardware



- **Next Steps**

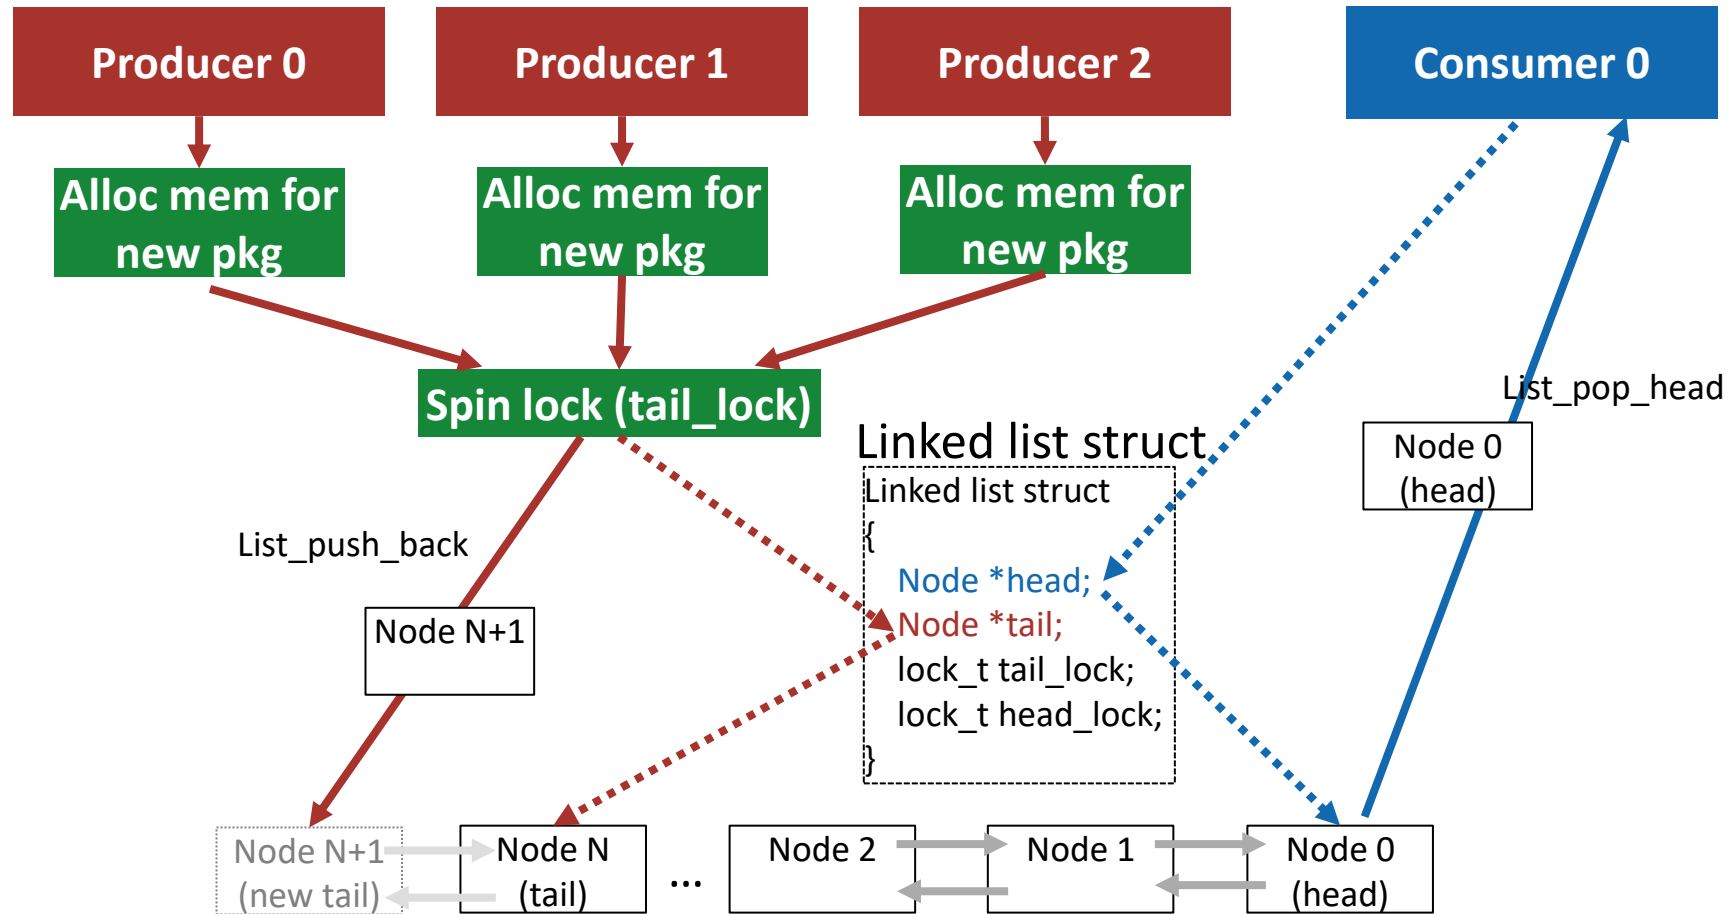
- Clean up the code and push to GitHub
- Currently use a single channel fake main memory in testbench => dramsys
- **Atomic support (high priority)**



Status Update: Software Plan



1. Multiple producers, single consumer double linked-list kernel




Status Update: Software Plan



2. Memory management runtime

- Assuming a fixed memory block size (a page) for each allocation and deallocation

1) Initial allocation: allocate new page from the sequential memory space



2) Free stage: push ptr of the deallocated page into a free mem linked list

3) Common allocation: pop the head node from the mem linked list, get the ptr to the new page



Status Update: Software



- **Progress of work**

1. Multiple producers, single consumer double linked-list kernel

- Multiple producers, single consumer read/write shared linked list
- TODO:
 - The simulation of the data movement of the consumer
 - Try to use multiple cores for data movement
 - Full atomic operation support from HW

2. Memory management runtime

- Dynamic memory allocation, deallocation, and reclamation at fine granularity (per package)
- TODO:
 - Variable allocation granularity



Thank you!

Q&A

