

# CachePool: Many-core cluster of customizable, lightweight scalar-vector PEs for irregular L2 data-plane workloads

Integrated Systems Laboratory (ETH Zürich)

**Zexin Fu, Diyou Shen**

zexifu, dishen@iis.ee.ethz.ch

**Alessandro Vanelli-Coralli**

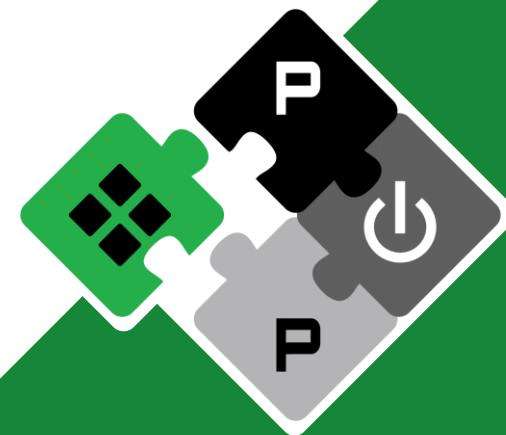
avanelli@iis.ee.ethz.ch

**Luca Benini**

lbenini@iis.ee.ethz.ch

**PULP Platform**

Open Source Hardware, the way it should be!



@pulp\_platform



pulp-platform.org



youtube.com/pulp\_platform



# RLC entity data structure



```
struct RLCEntity {
    unsigned int rlcId;
    unsigned int cellId; /* Indicates the cell to which the RLC entity belongs. */
    unsigned pollPdu; -- 32
    unsigned pollByte; -- 25000
    unsigned pduWithoutPoll; -- 0 (initial value) /* Indicates the total number of PDUs that are not polled. */
    unsigned byteWithoutPoll; -- 0 (initial value) /* Indicates the total bytes of PDUs that are not polled. */

    unsigned int sduNum; /* Number of sdus to be sent */ --- 3 -> 3+1
    unsigned int sduBytes; /* Number of sdus bytes to be sent */ --- 3*800 -> 4*800 (Assume that the length of each SDU is 800 bytes.)
    void *sduLinkHdr; /* First SDU to be sent */
    void *sduLinkTail; /* Last SDU to be sent */
    unsigned int Reserve1[64]** // Pieced together into a cacheline

    unsigned int vtNextAck; /* First SN to be confirmed */ -- 0 (initial value)
    unsigned int vtNext; /* Next Available RLCSN */ --- 0 (initial value)

    unsigned int sendPduNum; /* Number of pdus to be confirmed */ --- 0 (initial value)
    unsigned int sendPduBytes; /* Number of pdus to be confirmed */ --- 0 (initial value)
    void *waitAckLinkHdr; /* First SDU to be confirmed */
    void *waitAckLinkTail; /* Last SDU to be confirmed */
    unsigned int Reserve2[64]** // Pieced together into a cacheline
};
```

- **RLCEntity: 16 \* 4 byte variables + 2 \* reserved field**
  - Q1: How large should each of the reserved field be? 256 byte - ?
  - Q2: Should it be aligned to cache line?
  - Q3: Should we always load it to cache when the core access the entity?
  - In the following estimation we consider the entity size as  $16*4+2*256 = 576$  byte



# RLC Packet Handling



- **Test Case 1**

- **DownLink**

- RLC receives **one** linked-list of **1157** nodes in a **1/2000** second slot from PDCP
    - Throughput: **20Gbps**
    - Received packet payload size: **1350 Byte**
    - **Active RLC entity data structure size: 576 Byte**

- **UpLink**

- RLC receives **one** linked-list of **1953** nodes in a **1/2000** second slot from UE
    - Throughput: **1Gbps**
    - Received packet payload size: **160 Byte**



# RLC Packet Handling



- **Test Case 2**

- **DownLink**

- RLC receives **48** linked-lists **x 101** nodes (4882 nodes), in a **1/2000** second slot from PDCP
    - Throughput: **50Gbps**
    - Received packet payload size: **800 Byte**
    - **Active RLC entity data structure size: 27 KB**

- **UpLink**

- RLC receives **48** linked-list **x 325** nodes (15625 nodes) in a **1/2000** second slot from UE
    - Input packet payload size: **160 Byte**



# RLC Packet Handling



- **Test Case 3**

- **DownLink**

- RLC receives **256** linked-lists **x 8** nodes (1963 nodes), in a **1/2000** second slot from PDCP
    - Throughput: **20Gbps**
    - Received packet payload size: **800 Byte**
    - **Active User: 4800**
    - **Active RLC entity data structure size per slot: 144 KB**

- **UpLink**

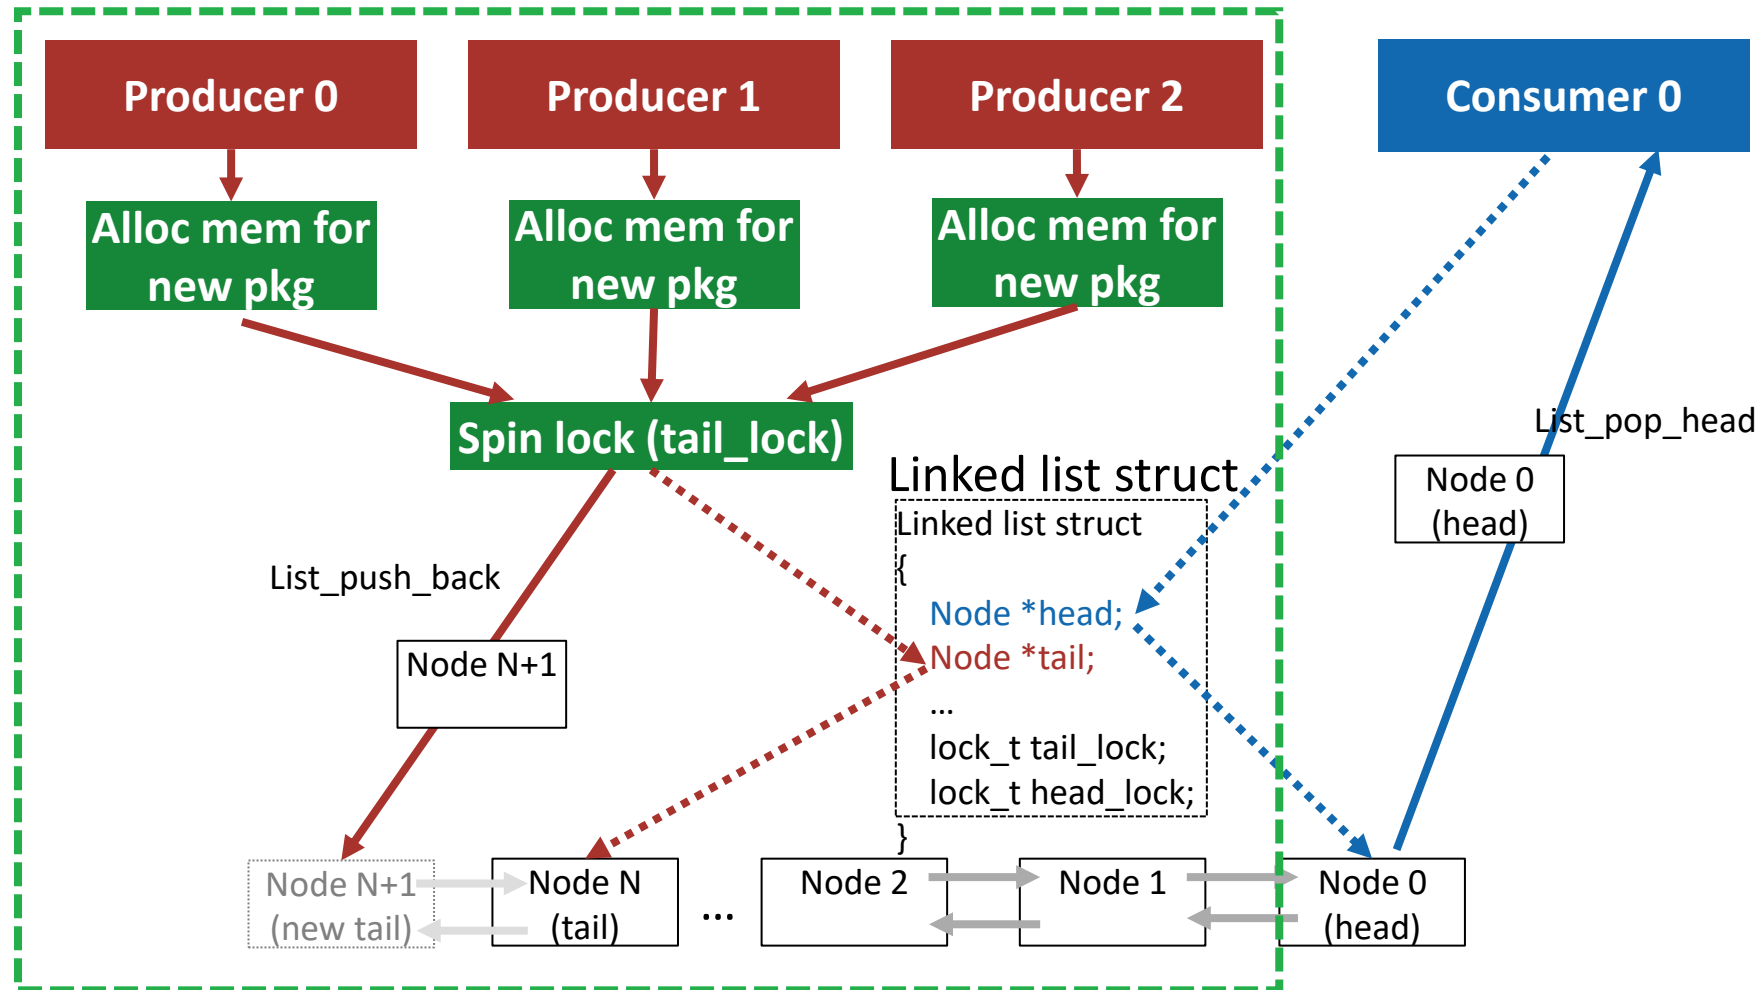
- RLC receives **384** linked-list **x 20** nodes (7813 nodes) in a **1/2000** second slot from UE
    - Input packet payload size: **160 Byte**
    - **Active User: 4800**



# Recap: Multi-producers, single consumer linked list



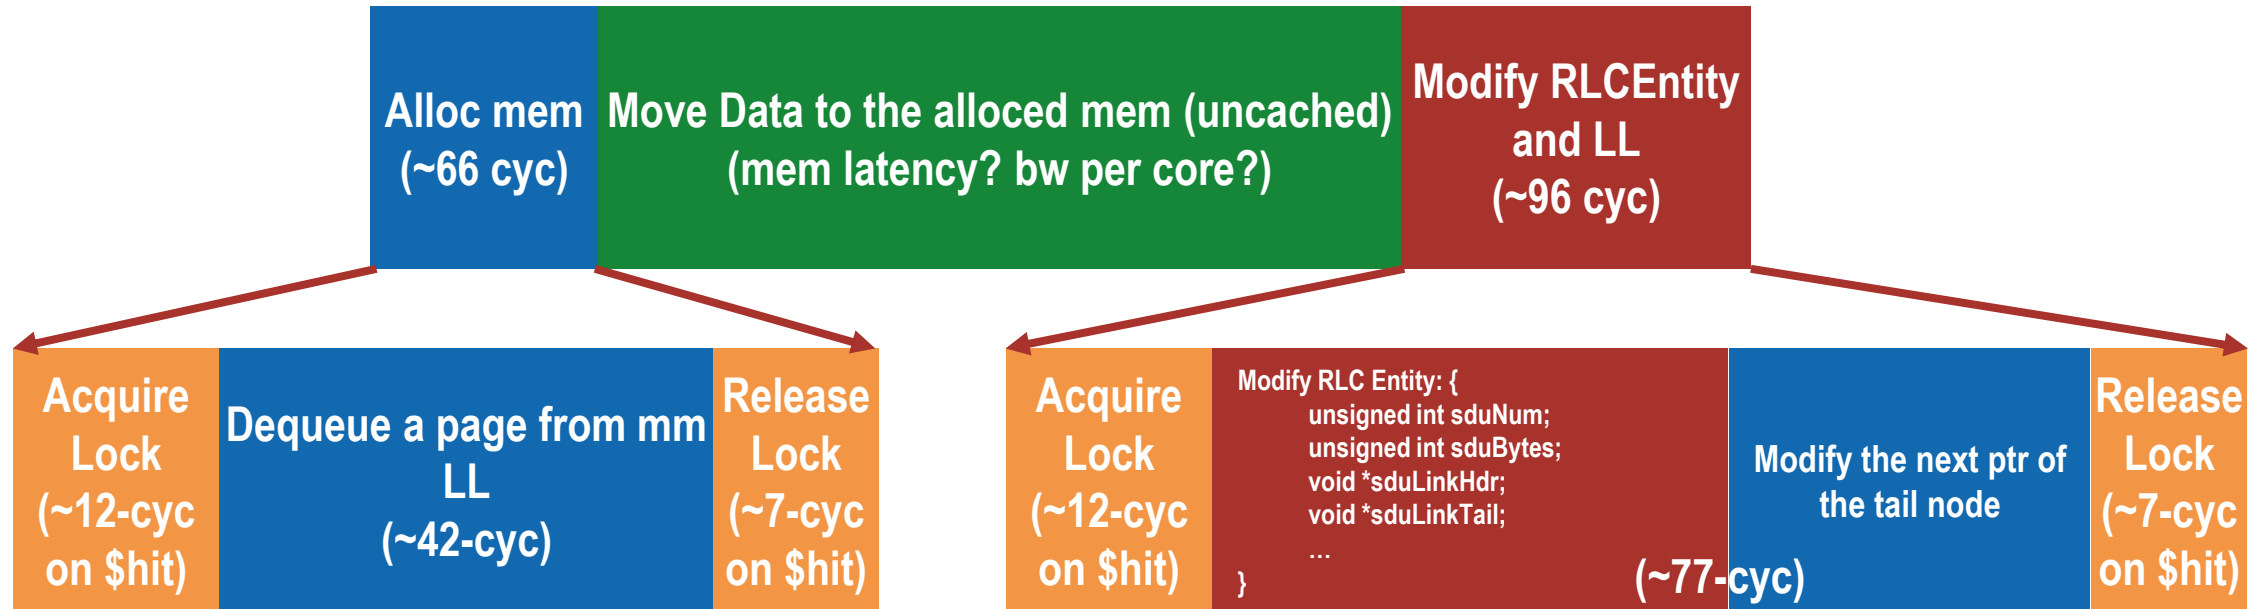
## 1. Multiple producers, single consumer double linked-list kernel



# Preliminary Estimation: Timing Breakdown



- **Packet Receiving Processing Task (Producer)**



- To move a 1350 byte data:

- If use HBM, e.g. latency 150-cyc, BW per core 256-bit/cyc:  $\sim 150 + 1.3\text{KB} / (256\text{bit/cyc}) = 192$  cycles, 354 cyc/pkg
- If use DRAM, e.g. latency 50-cyc, BW per core 64-bit/cyc:  $\sim 50 + 1.3\text{KB} / (64\text{bit/cyc}) = 218$  cycles, 380cyc/pkg

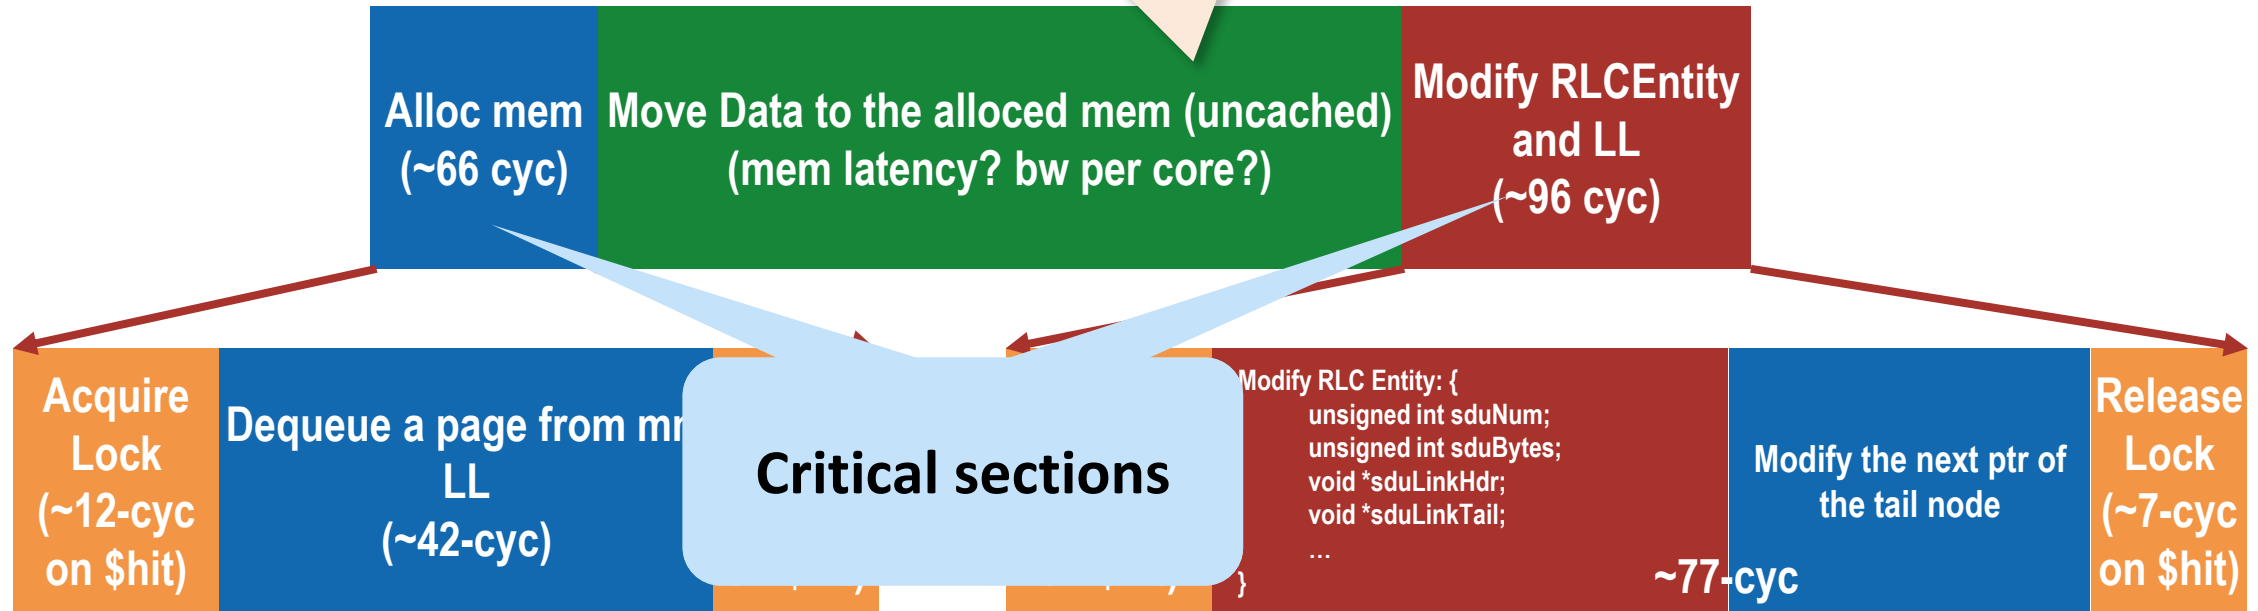


# Preliminary Estimation: Timing Breakdown



- **Packet Receiving Processing Task**

Parallelizable



- To move a 1350 byte data:

- If use HBM, e.g. latency 150-cyc, BW per core 256-bit/cyc:  $\sim 150 + 1.3\text{KB} / (256\text{bit/cyc}) = 192$  cycles, 354 cyc/pkg
- If use DRAM, e.g. latency 50-cyc, BW per core 64-bit/cyc:  $\sim 50 + 1.3\text{KB} / (64\text{bit/cyc}) = 218$  cycles, 380cyc/pkg






# Preliminary Estimation: In multicore scenario



- **Packet Receiving Processing Task**



- If use a tile (4 PEs) to process receiving task, run at 800MHz clk frequency
  - ~480 cycles to receive a pkg per core
    - **TODO:** The “pkg head processing” and the “insertion of sent pkg linked-list” haven’t been analyzed
  - With above preliminary estimation, each tile can receive ~3300 nodes per DL slot
    - Can handle test case 1;
    - Test case 2: may need 2 tiles;
    - Test case 3: the cache size of one tile need to be larger, e.g. 256 KB
    - Note: Other tasks (e.g. consumer, linear algebra ...) are not considered here.



# Status Update



- **Implement and analyze the “Packet Receiving Processing Task “**
  - Multi-user pkg receiving using a tile
  - Analyze the timing breakdown of the pkg receiving process
  - Fix several bug of the kernel
- **Question**
  - How should we add RLC header to the received pkg?
- **Next**
  - The pkg RLC header processing
  - The sent pkg linked-list insertion and deletion and its data movement
  - The ack message from the UE



# Thank you!

## Q&A

