

CachePool: Many-core cluster of customizable, lightweight scalar-vector PEs for irregular L2 data-plane workloads

Integrated Systems Laboratory (ETH Zürich)

Zexin Fu, Diyou Shen

zexifu, dishen@iis.ee.ethz.ch

Alessandro Vanelli-Coralli


avanelli@iis.ee.ethz.ch

Luca Benini

lbenini@iis.ee.ethz.ch

PULP Platform

Open Source Hardware, the way it should be!



@pulp_platform



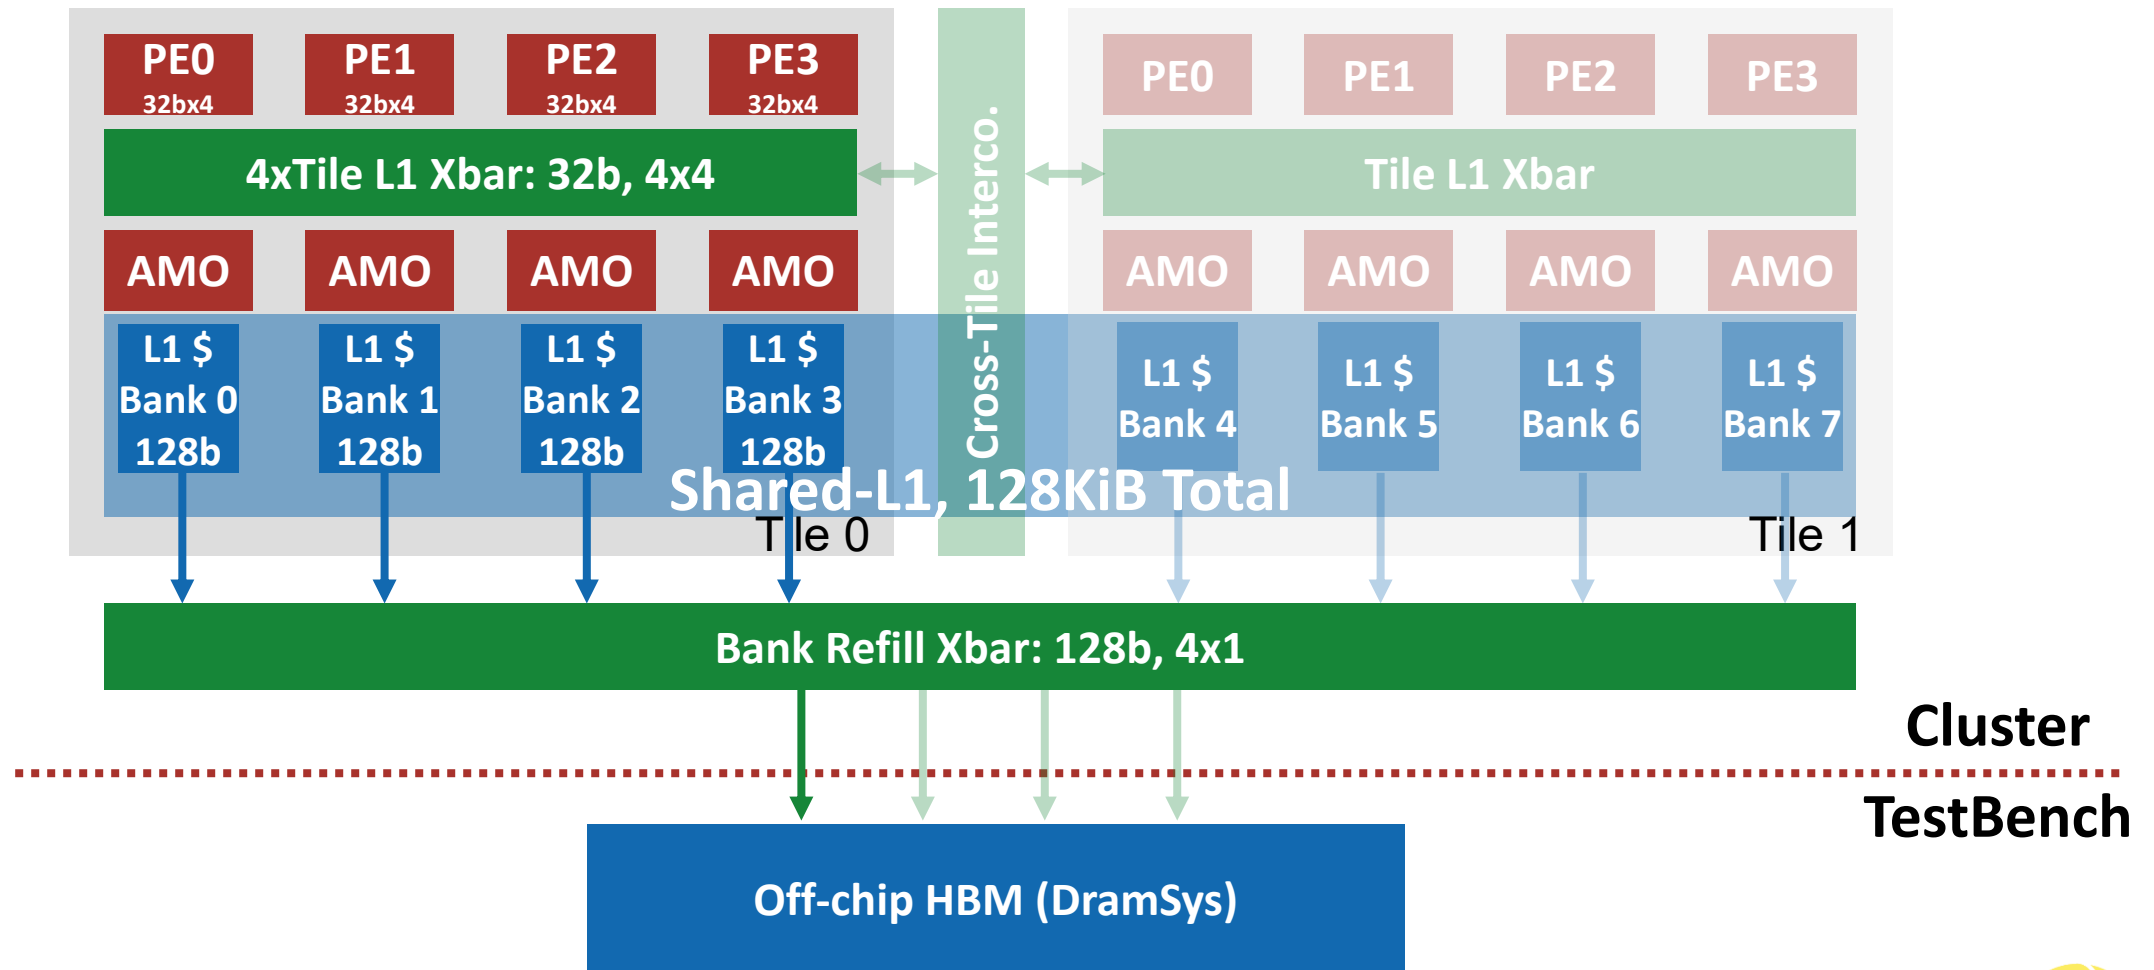
pulp-platform.org



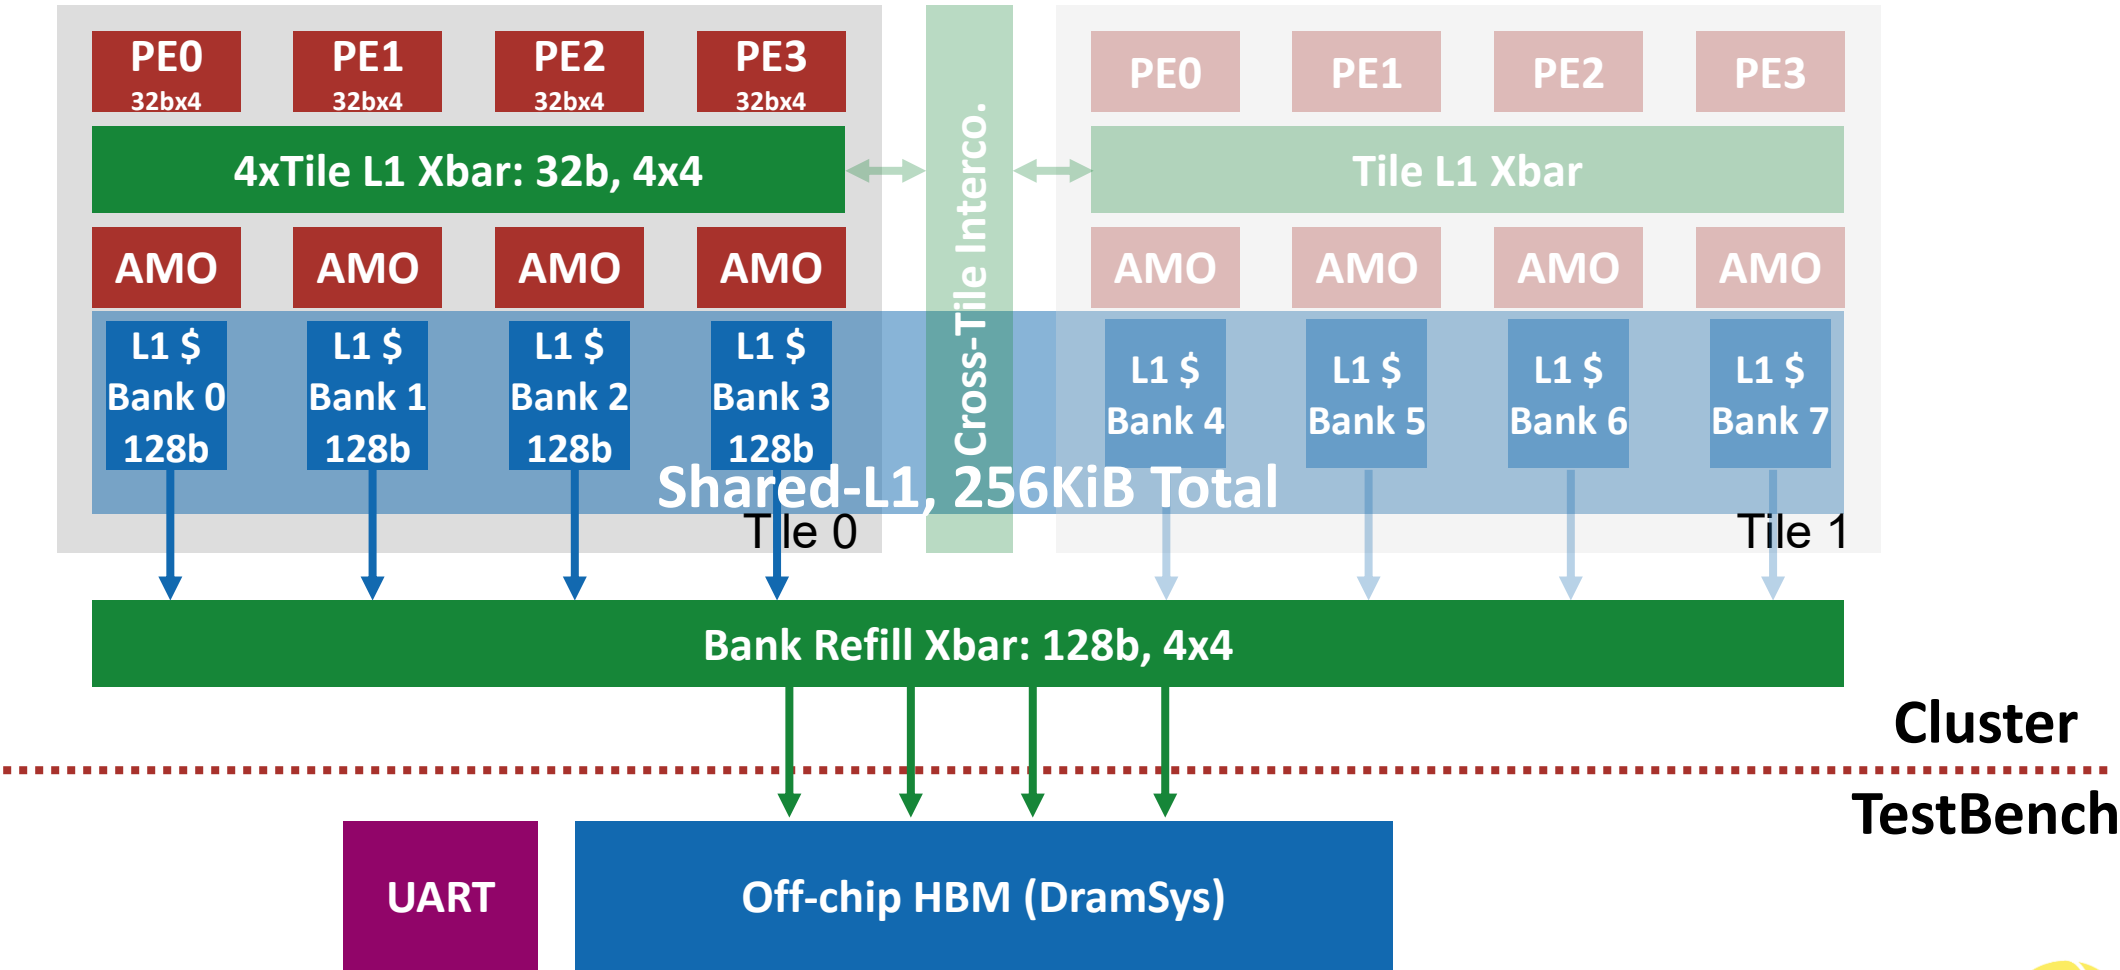
youtube.com/pulp_platform



Hardware Development



Hardware Development



Hardware Development



- **Complete**

- Integration of DramSys Simulator (100%)
- Cache-Refill Xbar
- UART module and printing support

- **In Progress**

- Code Cleanup: Remove the remaining modules from old cluster
 - DMA, SPM interconnection, stack modification
- PPA Analysis: Start with the existing 12nm flow
 - Faster setup
 - Easier area comparison with previously collected data
- Transit to 7nm flow for delivery

- **TODO**

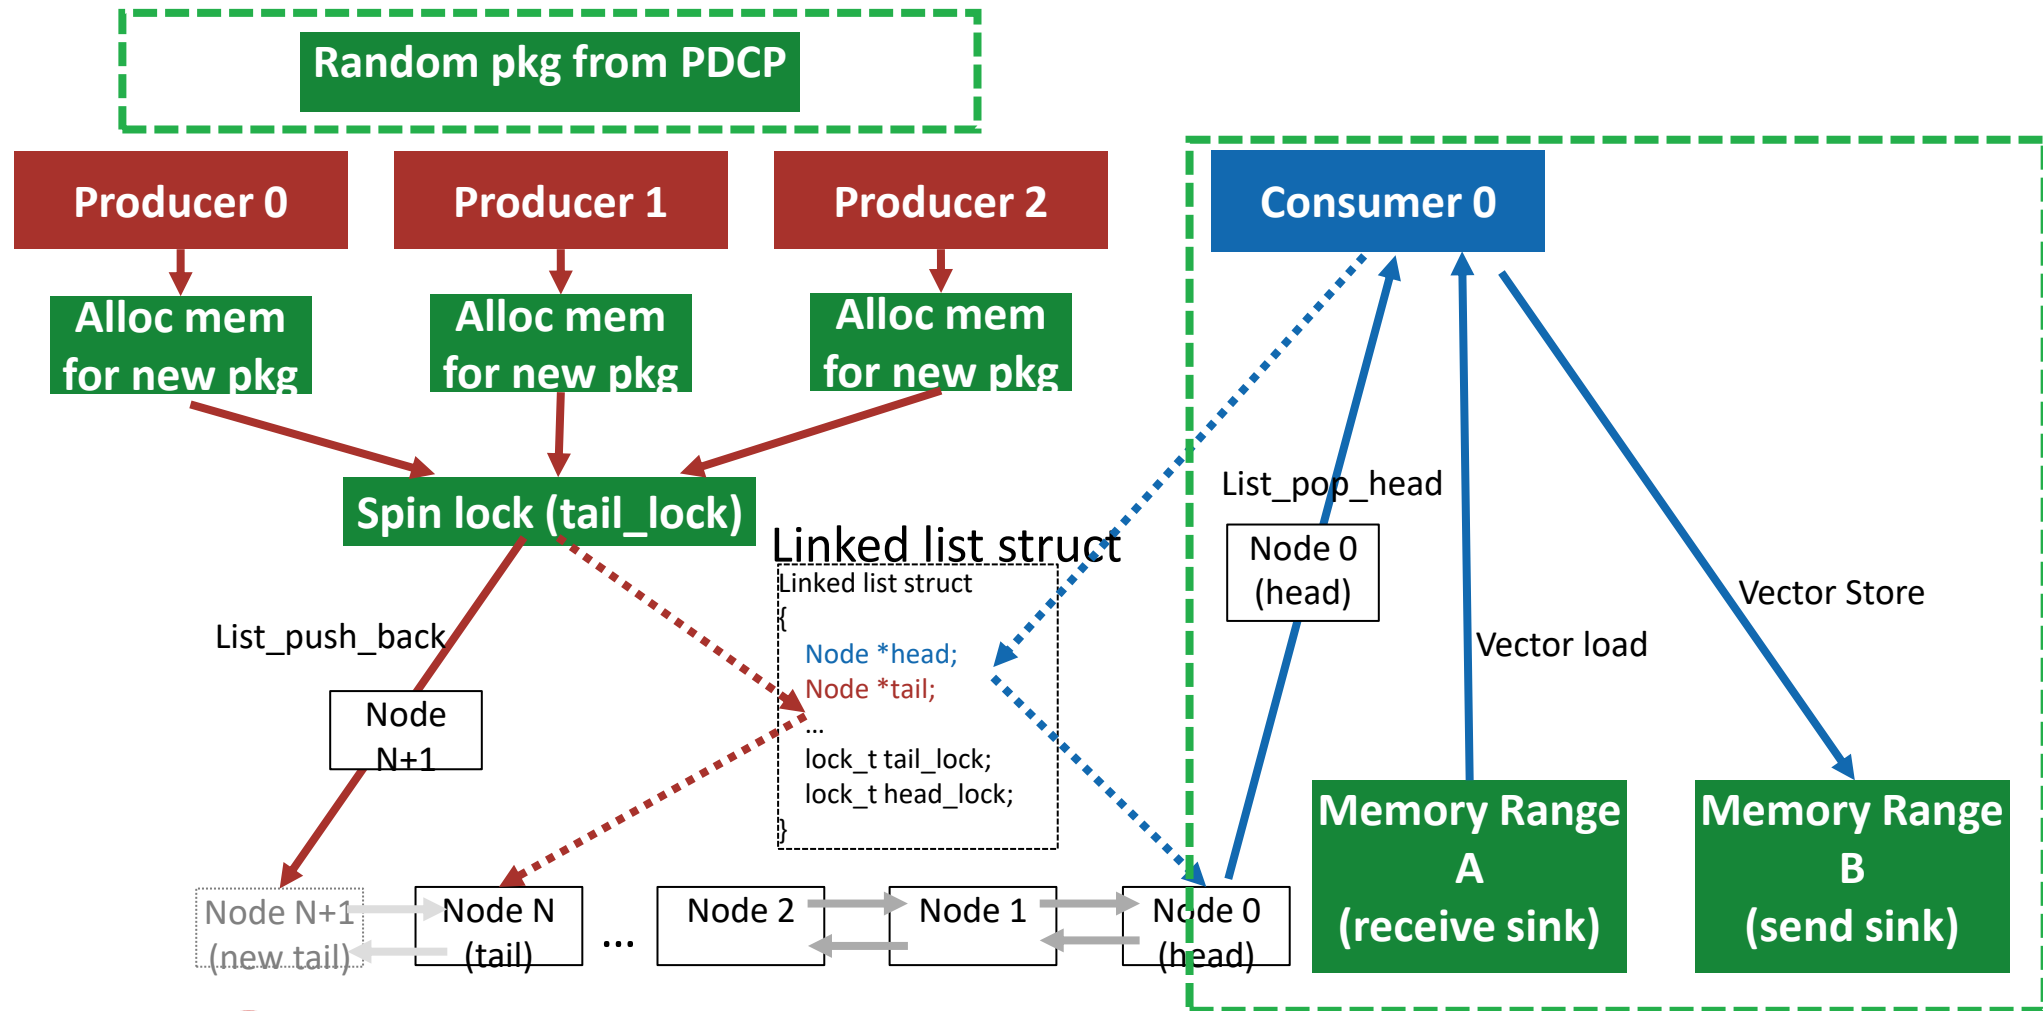
- Partition support



Software Development



Use Spatz VLSU to move data, simulate the RLC pkg sending process





- ```
"active_user_number": 1,
"pkg_length": 1350,
"pdcp_header_length": 6,
"src_addr": "0xA0000000",
"src_length": 268435456, // 268435456 bytes = 256 MiB
"tgt_addr": "0xB0000000",
"tgt_length": 268435456, // 268435456 bytes = 256 MiB
"total_pkg_number": 1000 // max 198400 packages =
268435456 bytes / (1350 + 3) bytes per package
```

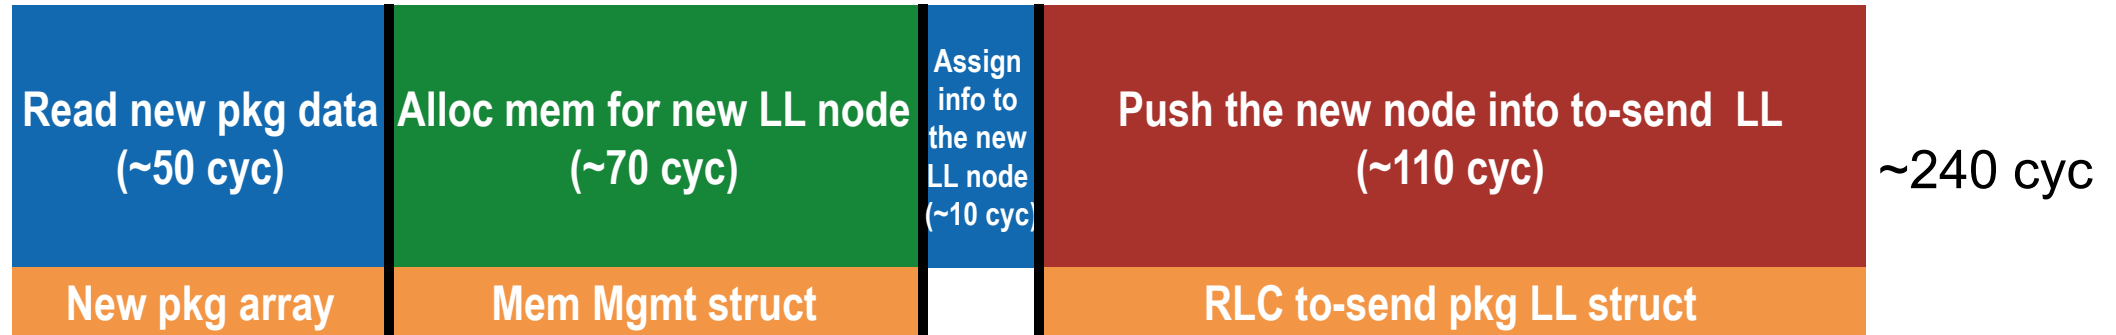
**ETH** zürich




# Packet Receiving Task Preliminary Timing Breakdown



- **Producer: Receive new pkg from PDCP layer, add its info into to-send LL**



- **Consumer: Pop the pkg info from the LL, move pkg data (transmit to UE)**



- Assume Freq 1 GHz
- BW 2142857pkg/s \* 2KB = 4GBps (for now not enough, challenge: data movement latency)





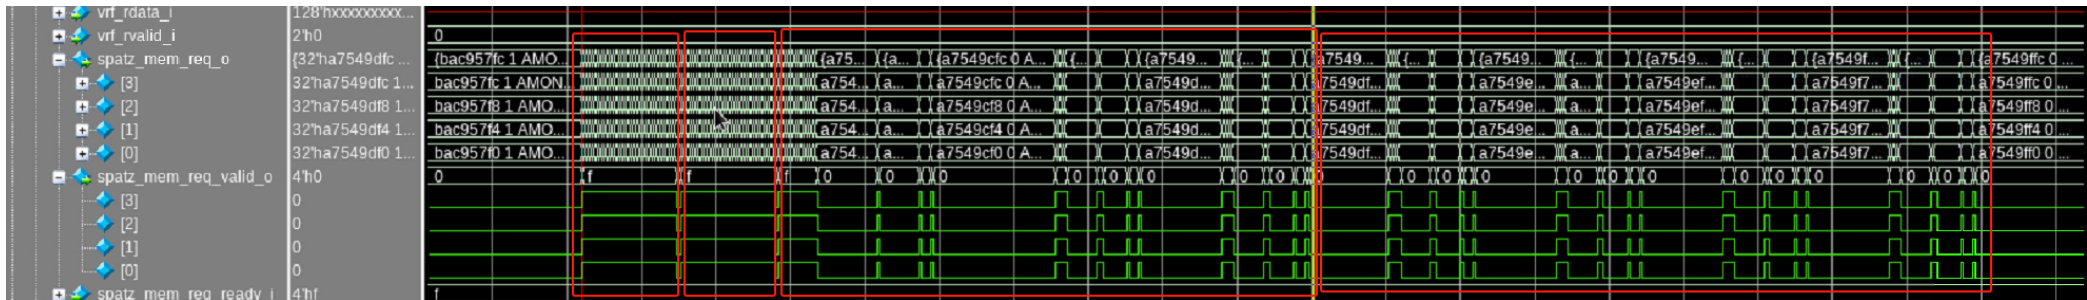
- 
- Figure 1: Scheduling diagram for four cores. The diagram shows the execution of tasks 1, 2, and 3 across four cores. Core 0 shows a sequence of tasks (1, 2, 3) in a repeating pattern. Core 1 starts with task 3, then a long blue bar, then task 2. Core 2 starts with task 3, then a long blue bar, then task 2. Core 3 starts with task 3, then a long blue bar, then task 2. The diagram illustrates the execution of tasks 1, 2, and 3 across four cores, with task 3 being the longest and task 2 being the shortest.

- ETH** zürich  ALMA MATER STUDIORUM  
UNIVERSITÀ DI BOLOGNA

# Memory Movement Challenge



- **Each Spatz can only achieve 2 byte/cycle (max bw is 16 byte/cycle)**
  - Bottleneck: The size of Spatz VLSU ROB can handle only 256 – 512 byte in-flight load
    - Once the ROB is full, the VLSU has to wait until the previous load data back.



- Possible solution
  - Redesign the ROB of the VLSU, make it more scalable and can handle more in-flight load
  - Offload the memory moving operations to an extra controller (prefetcher/dma/cache controller ...)



# Software Development



- **Complete**

- Random generated pkg from PDCP layer
- Design and evaluate the RLC data management kernel with VLSU for data movement

- **In Progress**

- GEMM and GEMV kernel performance calibration: HW is not stable for now
- The sent linked-list for the RLC data management kernel

- **TODO**

- Explore the way to fully utilize the memory BW for RLC pkg data movement



# Thank you!

## Q&A

